

The Knora Base Ontology

Lukas Rosenthaler, Benjamin Geer, Tobias Schweizer, Ivan Subotic
Digital Humanities Lab, University of Basel

26 January 2016

Todo list

Issue 290: Should this be in the triplestore?	3
Give more details on how this works.	11
Issue 292: Nonterminal nodes do not exist yet in the ontology.	15
How do we differentiate between truly hierarchical structures and structures that are just encoded as a hierarchy for a given XML file?	16

Contents

1 Introduction	1
1.1 Resource Description Framework (RDF)	1
2 Projects	3
3 The Knora Data Model	3
3.1 Resources	3
3.1.1 Properties of Resource	4
3.1.2 Representations	5
3.2 Values	6
3.2.1 Properties of Value	6
3.3 Subclasses of Value	7
3.3.1 TextValue	7
3.3.2 DateValue	7
3.3.3 IntValue	8
3.3.4 FloatValue	8
3.3.5 TimeValue	8
3.3.6 GeomValue	8

3.3.7	GeonameValue	8
3.3.8	IntervalValue	8
3.3.9	ListValue	9
3.3.10	FileValue	9
3.3.11	LinkValue	10
3.3.12	ExternalResValue	11
4	Links Between Resources	11
5	Text with Standoff Nodes	14
5.1	Subclasses of Standoff	17
5.1.1	StandoffVisualAttribute	17
5.1.2	StandoffHref	17
5.1.3	StandoffLink	17
6	Authorization	18
6.1	Users and Groups	18
6.2	Permissions	20
7	Open Questions	21
7.1	Extending Existing Resource Definitions	21

1 Introduction

1.1 Resource Description Framework (RDF)

Knora¹ uses a hierarchy of ontologies based on RDF [8], RDF Schema (RDFS)[2], and Web Ontology Language (OWL)[6]. Both RDFS and OWL are expressed in RDF. RDF expresses information as a set of statements (called *triples*). A triple consists of a subject, a predicate, and an object (Figure 1).




Figure 1: An RDF triple.

The object may be either a literal value (such as a name or number) or another subject. Thus it is possible to create complex graphs that connect many subjects, as in Figure 2.

¹**Knowledge Organization, Representation, and Annotation**, <http://knora.org>.

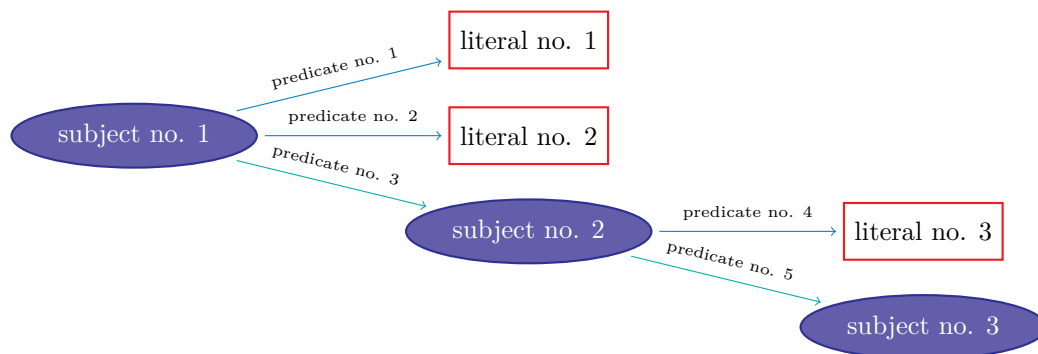


Figure 2: An RDF graph.

RDF uses unique, URL-like identifiers called Internationalized Resource Identifiers (IRIs)[5]. Anything that can be a subject has an IRI, as does each predicate. Within a given project, IRIs typically differ only in their last component (the ‘local part’), which is often the fragment following a # character. Such IRIs share a long ‘prefix’. In Turtle[1] and similar formats for writing RDF, a short prefix label can be defined to represent the long prefix. Then an IRI can be written as a prefix label and a local part, separated by a colon ‘:’. For example, if the ‘example’ project’s long prefix is `http://www.example.org/rdf#`, and it contains subjects with IRIs like `http://www.example.org/rdf#book`, we can define the prefix label ‘ex’ to represent the prefix label, and write prefixed names for IRIs as in Figure 3. In this document, we use the prefix label **kb** to represent the Knora base ontology,² but we usually omit it for brevity.

2 Projects

In Knora, each item of data belongs to some particular project. Each project using Knora must define a **kb:knoraProject**, which has these properties (cardinalities are indicated in parentheses after each property name):

shortname (1) A short name that can be used to identify the project in configuration files and the like.

basepath (1) The filesystem path of the directory where the project’s files are stored.

foaf:name (0-1) The name of the project.

²<http://www.knora.org/ontology/knora-base#>

Issue 290:
Should this
be in the
triplestore?

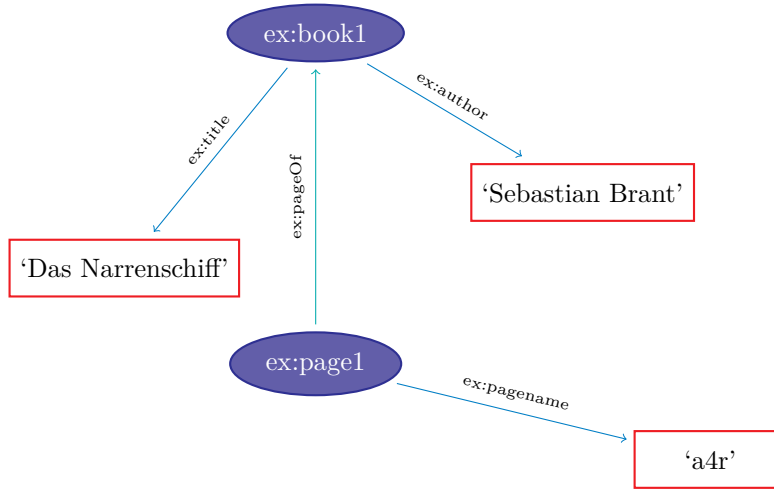


Figure 3: An RDF graph written with prefix labels.

description (0-1) A description of the project.

belongsTo (0-1) The `kb:Institution` that the project belongs to.

Resources and values are associated with a project by means of the `kb:attachedToProject` property, as described in Section 3. Users are associated with a project by means of the `kb:isInProject` property, as described in Section 6.

3 The Knora Data Model

The Knora data model is based on the observation that, in the humanities, a value or literal is often itself structured and can be highly complex. Moreover, a value may have its own metadata, such as its creation date, information about ownership, permissions, and so on. Therefore, the Knora base ontology describes structured value types that can store this type of metadata. For example, in Figure 4, a book (`ex:book2`) has a title (identified by the predicate `ex:title`) and a publication date (`ex:pubdate`), each of which has some metadata.

3.1 Resources

Each project using the Knora base ontology defines its own OWL classes, derived from `kb:Resource`, to represent the types of data it deals with. Each

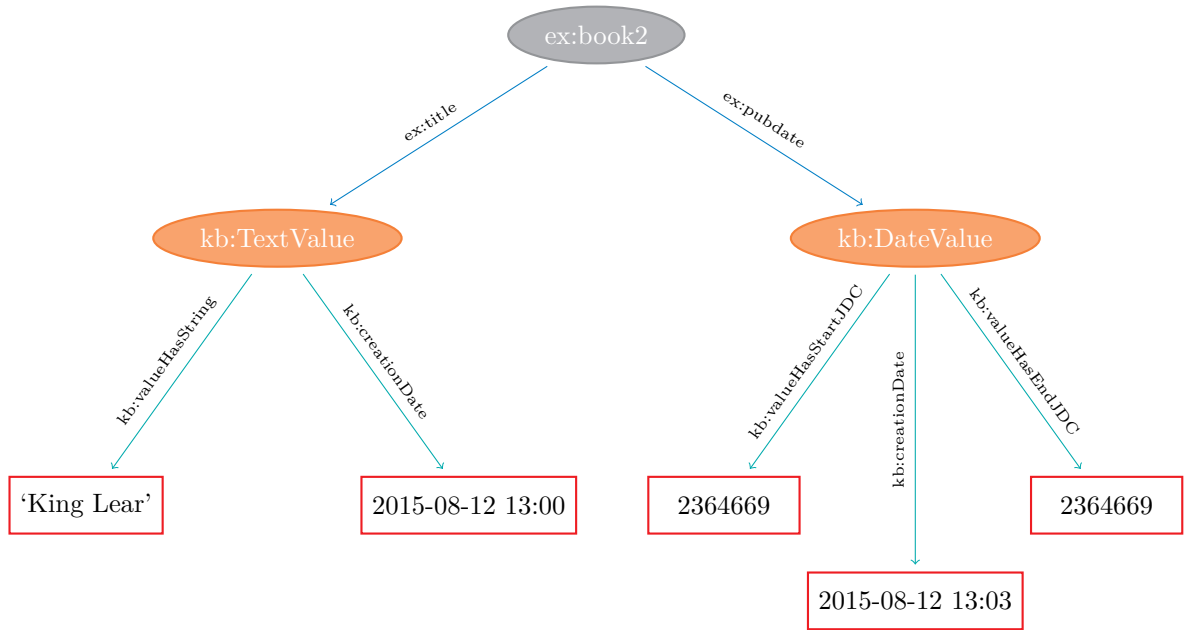


Figure 4: Structured values.

project-specific resource class definition must use OWL cardinality restrictions to specify the properties that a given type of resource can have. Properties that point to Knora values (see Section 3.2) must be derived from `kb:hasValue`, and properties that point to other resources must be derived from `kb:hasLinkTo`. A few such properties are defined in the Knora base ontology, but projects typically need to define many of their own properties.

Resources are not versioned; only their values are versioned (see Section 3.2). A resource can be marked as deleted. It then remains in the repository, but becomes invisible to ordinary queries. It cannot be undeleted, because even though resources are not versioned, it is necessary to be able to find out when a resource was deleted. If desired, a new resource can be created by copying data from a deleted resource.

3.1.1 Properties of Resource

creationDate (1) The time when the resource was created.

attachedToUser (1) The user who owns the resource.

attachedToProject (1) The project that the resource is part of.

lastModificationDate (0-1) A timestamp indicating when the resource (or one of its values) was last modified.

seqnum (0-1) The sequence number of the resource, if it is part of an ordered group of resources, such as the pages in a book.

isDeleted (0-1) Indicates whether the resource has been deleted.

deleteDate (0-1) If the resource has been deleted, indicates when it was deleted.

Resources can have properties that point to other resources; see Section 4. A resource grants permissions to groups of users; see Section 6.

3.1.2 Representations

It is not practical to store all data in RDF. In particular, RDF is not a good storage medium for binary data such as images. Therefore, Knora stores such data outside the triplestore, in ordinary files. A resource can have one or more files attached to it. For each file, there is a `kb:FileValue` in the triplestore containing metadata about the file (see Section 3.3.10). A resource that has file values must belong to one of the subclasses of `kb:Representation`. The base class `Representation`, which is not intended to be used directly, has this property:

hasFileValue (1-n) Points to one or more file values.

Its subclasses, which are intended to be used directly in data, include:

StillImageRepresentation A representation containing still image files.

MovingImageRepresentation A representation containing video files.

AudioRepresentation A representation containing audio files.

DDDrepresentation A representation containing 3D images.

TextRepresentation A representation containing formatted text files, such as XML files.

DocumentRepresentation A representation containing documents (such as PDF files) that are not text files.

There are two ways for a project to design classes for representations. The simpler way is to create a resource class that represents a thing in the world (such as `ex:Painting`) and also belongs to a subclass of `Representation`. This is adequate if the class can have only one type of file attached to it. For

example, if paintings are represented only by still images, `ex:Painting` could be a subclass of `StillImageRepresentation`. This is the only approach supported in version 1 of the Knora API.

The more flexible approach, which is allowed by the Knora base ontology and will be supported by version 2 of the Knora API, is for each `ex:Painting` to use the `kb:hasRepresentation` property to point to other resources containing files that represent the painting. Each of these other resources can extend a different subclass of `Representation`. For example, a painting could have a `StillImageRepresentation` as well as a `DDDrepresentation`.

3.2 Values

The Knora base ontology defines a set of OWL classes that are derived from `kb:Value` and represent different types of structured values found in humanities data. This set of classes may not be extended by project-specific ontologies.

A value is always part of one particular resource, which points to it using some property derived from `hasValue`. For example, a project-specific ontology could specify a `Book` class with a property `hasSummary` (derived from `hasValue`), and that property could have an `rdfs:range` of `TextValue`. This would mean that the summary of each book is represented as a `TextValue`.

Knora values are versioned. Existing values are not modified. Instead, a new version of an existing value is created. The new version is linked to the old version via the `previousValue` property.

‘Deleting’ a value means creating a new version, marked with `kb:isDeleted` and pointing to the previous version. A triple then points from the containing resource to this new, deleted version of the value, which remains in the repository but is invisible to ordinary queries. To simplify the enforcement of ontology constraints, and for consistency with resource updates, no new versions of a deleted value can be made; it is not possible to undelete. Instead, if desired, a new value can be created by copying data from a deleted value.

3.2.1 Properties of Value

valueCreationDate (1) The date and time when the value was created.

attachedToUser (1) The user who owns the value.

attachedToProject (0-1) The project that the value is part of. If not specified, defaults to the project of the containing resource.

valueHasString (1) A human-readable string representation of the value's contents, which is available to Knora's full-text search index.

valueHasOrder (0-1) A resource may have several properties of the same type with different values (which will be of the same class), and it may be necessary to indicate an order in which these values occur. For example, a book may have several authors which should appear in a defined order. Hence, **valueHasOrder**, when present, points to an integer literal indicating the order of a given value relative to the other values of the same property. These integers will not necessarily start at any particular number, and will not necessarily be consecutive.

previousValue (0-1) The previous version of the value.

isDeleted (0-1) Indicates whether the value has been deleted.

Each Knora value can grant permissions (see Section 6).

3.3 Subclasses of Value

3.3.1 TextValue

Represents text, possibly including markup. Property:

valueHasStandoff (0-n) Points to a Standoff node. See Section 5.

3.3.2 DateValue

Humanities data includes many different types of dates. In Knora, a date has a specified calendar, and is always represented as a period with start and end points (which may be equal), each of which has a precision (**DAY**, **MONTH**, or **YEAR**). Internally, the start and end points are stored as two Julian Day Numbers. This calendar-independent representation makes it possible to compare and search for dates regardless of the calendar in which they were entered. Properties:

valueHasCalendar (1) The name of the calendar in which the date should be displayed. Currently **GREGORIAN** and **JULIAN** are supported.

valueHasStartJDC (1) The Julian Day Number of the start of the period (an `xsd:integer`).

valueHasStartPrecision (1) The precision of the start of the period.

valueHasEndJDC (1) The Julian Day Number of the end of the period
(an `xsd:integer`).

valueHasEndPrecision (1) The precision of the end of the period.

3.3.3 IntValue

Represents an integer. Property:

valueHasInteger (1) An `xsd:integer`.

3.3.4 FloatValue

Represents a floating-point number. Property:

valueHasFloat (1) An `xsd:float`.

3.3.5 TimeValue

Represents a precise point on a timeline, e.g. relative to the beginning of an audio or video file. Property:

valueHasTime (1) An `xsd:float` representing a number of seconds from the beginning of the timeline.

3.3.6 GeomValue

Represents a geometrical object as a JSON string, using normalized coordinates. Property:

valueHasGeometry (1) A JSON string.

3.3.7 GeonameValue

Represents a geolocation, using the numerical codes found at <http://geonames.org>. Property:

valueHasGeonameCode (1) the numerical code of a geographical feature from <http://geonames.org>, represented as an `xsd:integer`.

3.3.8 IntervalValue

Represents a time interval, with precise start and end times on a timeline, e.g. relative to the beginning of an audio or video file. Properties:

valueHasIntervalStart (1) An `xsd:float` representing the start of the interval in seconds.

valueHasIntervalEnd (1) An `xsd:float` representing the end of the interval in seconds.

3.3.9 ListValue

Projects often need to define lists or hierarchies of categories that can be assigned to many different resources. Then, for example, a user interface can provide a drop-down menu to allow the user to assign a category to a resource. The `ListValue` class provides a way to represent these sorts of data structures. It can represent either a flat list or a tree.

A `ListValue` has this property:

valueHasListNode (1) Points to the root `ListNode` of the list or tree.

Each `ListNode` can have the following properties:

hasSubListNode (0-n) Points to the node's child nodes, if any.

listNodePosition (1) An integer indicating the node's position in the list of its siblings.

isRootNode (0-1) Set to `true` if this is the root node.

listNodeName (0-n) The node's human-readable name.

3.3.10 FileValue

Knora stores certain kinds of data outside the triplestore, in files (see Section 3.1.2). Each digital object that is stored outside the triplestore has associated metadata, which is stored in the triplestore in a `kb:FileValue`. The base class `FileValue`, which is not intended to be used directly, has these properties:

internalFilename (1) The name of the file as stored by the Knora API server.

internalMimeType (1) The MIME type of the file as stored by the Knora API server.

originalFilename (0-1) The original name of the file when it was uploaded to the Knora API server.

originalMimeType (0-1) The original MIME type of the file when it was uploaded to the Knora API server.

isPreview (0-1) A boolean indicating whether the file is a preview, i.e. a small image representing the contents of the file. A preview is always a **StillImageFileValue**, regardless of the type of the enclosing **Representation**.

The subclasses of **FileValue**, which are intended to be used directly in data, include:

StillImageFileValue Contains metadata about a still image file.

MovingImageFileValue Contains metadata about a video file.

AudioFileValue Contains metadata about an audio file.

DDDFFileValue Contains metadata about a 3D image file.

TextFileValue Contains metadata about a text file.

DocumentFileValue Contains metadata about a document (such as PDF) that is not a text file.

Each of these classes contains properties that are specific to the type of file it describes. For example, still image files have dimensions, video files have frame rates, and so on.

The files in a given representation must be semantically equivalent, meaning that coordinates that relate to one file must also be valid for other files in the same representation. Coordinates in Knora are expressed as fractions of the size of the object on some dimension; for example, image coordinates are expressed as fractions of its width and height, rather than in pixels. Therefore, the image files in a **StillImageRepresentation** must have the same aspect ratio, but they need not have the same dimensions in pixels. Similarly, the audio and video files in an **AudioRepresentation** or **MovingImageRepresentation** must have the same length in seconds, but may have different bitrates.

FileValue objects are versioned like other values, and the actual files stored by Knora are also versioned. Version 1 of the Knora API does not provide a way to retrieve a previous version of a file, but this feature will be added in a subsequent version of the API.

3.3.11 LinkValue

A **LinkValue** is an RDF ‘reification’ containing metadata about a link between two resources. It is therefore a subclass of **rdf:Statement** as well as of **Value**. It has these properties:

rdf:subject (1) The resource that is the source of the link.

rdf:predicate (1) The link property.

rdf:object (1) The resource that is the target of the link.

valueHasRefCount (1) The reference count of the link. This is meaningful when the **LinkValue** describes resource references in Standoff text markup (see Section 5.1.3). Otherwise, the reference count will always be 1 (if the link exists) or 0 (if it has been deleted).

For details about how links are created in Knora, see Section 4.

3.3.12 ExternalResValue

Represents a resource that is not stored in the RDF triplestore managed by the Knora API server, but instead resides in an external repository managed by some other software. The **ExternalResValue** contains the information that the Knora API server needs in order to access the resource, assuming that a suitable gateway plugin is installed.

extResAccessInfo (1) The location of the repository containing the external resource (e.g. its URL).

extResId (1) The repository-specific ID of the external resource.

extResProvider (1) The name of the external provider of the resource.

Give more details on how this works.

4 Links Between Resources

A link between two resources is expressed, first of all, as a triple, in which the subject is the resource that is the source of the link, the predicate is a ‘link property’ (a subproperty of `kb:hasLinkTo`), and the object is the resource that is the target of the link.

It is also useful to store metadata about links. For example, Knora needs to know who owns the link, who has permission to modify it, when it was created, and so on. Such metadata cannot simply describe the link property, because then it would refer to that property in general, not to any particular instance in which that property is used to connect two particular resources. To attach metadata to a specific link in RDF, it is necessary to create an RDF ‘reification’. A reification makes statements about a particular triple (subject, predicate, object), in this case the triple that expresses the link between the resources. Knora uses reifications of type `kb:LinkValue` (described in Section 3.3.11) to store metadata about links.

For example, suppose a project describes paintings that belong to collections. The project can define an ontology as follows (expressed here in Turtle format, and simplified for the purposes of illustration):

```
@prefix kb <http://www.knora.org/ontology/knora-base#> .
@prefix : <http://www.knora.org/ontology/paintings#> .

:Painting rdf:type owl:Class ;
  rdfs:subClassOf kb:Resource ,
    [ rdf:type owl:Restriction ;
      owl:onProperty :hasName ;
      owl:cardinality 1 ] ,
    [ rdf:type owl:Restriction ;
      owl:onProperty :title ;
      owl:cardinality 1 ] .

:Collection rdf:type owl:Class ;
  rdfs:subClassOf kb:Resource ,
    [ rdf:type owl:Restriction ;
      owl:onProperty :hasName ;
      owl:cardinality 1 ] .

:hasName rdf:type owl:ObjectProperty ;
  rdfs:label "Name of artist" ;
  rdfs:domain :Painting ;
```

```

    rdfs:range kb:TextValue .

:title rdf:type owl:ObjectProperty ;
    rdfs:label "Title of painting"
    rdfs:domain :Painting ;
    rdfs:range kb:TextValue .

:hasName rdf:type owl:ObjectProperty ;
    rdfs:label "Name of collection" ;
    rdfs:domain :Collection ;
    rdfs:range kb:TextValue .

```

To link the paintings to the collection, we must add a ‘link property’ to the ontology. In this case, the link property will point from a painting to the collection it belongs to. Every link property must be a subproperty of `hasLinkTo`.

```

:isInCollection rdf:type owl:ObjectProperty ;
    rdfs:subPropertyOf kb:hasLinkTo ;
    rdfs:domain :Painting ;
    rdfs:range :Collection .

```

We must then add a ‘link value property’, which will point from a painting to a `LinkValue` (described in Section 3.3.11), which will contain metadata about the link between the property and the collection. The name of the link value property is constructed using a simple naming convention: the word `Value` is appended to the name of the link property. In this case, since our link property is called `isInCollection`, the link value property must be called `ex:isOnPageValue`. Every link value property must be a subproperty of `kb:hasLinkToValue`.

```

:isInCollectionValue rdf:type owl:ObjectProperty ;
    rdfs:subPropertyOf kb:hasLinkToValue ;
    rdfs:domain :painting ;
    rdfs:range kb:LinkValue .

```

Given this ontology, we can create some RDF data describing a painting and a collection:

```

@prefix paintings <http://www.knora.org/ontology/paintings#> .
@prefix data <http://www.knora.org/ontology/paintings/data#> .

```

```

data:dali_4587 rdf:type paintings:Painting ;
  paintings:title data:value_A ;
  paintings:hasName data:value_B .

data:value_A rdf:type kb:TextValue ;
  kb:valueHasString "The Persistence of Memory" .

data:value_B rdf:type kb:TextValue ;
  kb:valueHasString "Salvador Dali" .

data:pompidou rdf:type paintings:Collection ;
  paintings:hasName data:value_C .

data:value_C rdf:type kb:TextValue ;
  kb:valueHasString "Centre Pompidou, Paris" .

```

We can then state that the painting is in the collection:

```

data:dali_4587 paintings:isInCollection data:pompidou ;
  paintings:isinCollectionValue data:value_D .

data:value_D rdf:type kb:LinkValue ;
  rdf:subject data:dali_4587 ;
  rdf:predicate paintings:isInCollection ;
  rdf:object data:pompidou ;
  kb:valueHasRefCount 1 ;

```

This creates a link (`isInCollection`) between the painting and the collection, along with a reification containing metadata about the link. We can visualise the result as the graph shown in Figure 5.

5 Text with Standoff Nodes

Knora is designed to be able to store text with markup, which can indicate formatting and structure, as well as the complex observations involved in transcribing handwritten manuscripts. One popular way of representing text in the humanities is to encode it as TEI/XML³ using the Text Encoding

³TEI refers both to an organization and an XML-based markup language (or more precisely: a set of grammar modules – XML schemas – that can be combined to define a markup language). For reasons of clarity, we use the term TEI/XML to refer to the markup language.




Figure 5: An RDF graph showing how two Knora resources are linked together.

Initiative (TEI) guidelines [3]. In Knora, a TEI/XML document can be stored as a file with attached metadata.

However, Knora also supports ‘standoff’ nodes, which are stored separately from the text. This has some advantages over embedded markup such as XML.⁴ While XML requires markup to have a hierarchical structure, and does not allow overlapping tags, standoff nodes do not have these limitations [7]. A ‘standoff node’ can assign an attribute to any substring in the text by giving its start and end positions.⁵ For example, suppose we have the following text:

This *sentence* **has overlapping visual** attributes.

This would require just two standoff nodes: (*italic*, start=5, end=29) and (**bold**, start=14, end=36).

⁴It is also possible to encode standoff markup using XML. For example, the TEI discusses standoff markup in its guidelines [3, chapters 16.9 and 20.4]. However, standoff markup is not widely applied in the TEI community. The main focus lies on encoding a hierarchy of elements.

⁵Unlike in corpus linguistics, we do not use any tokenization resulting in a form of predefined segmentation that would limit the user’s possibility to freely annotate any ranges in the text.

Moreover, standoff makes it possible to mark up the same text in different, possibly incompatible ways, allowing for different interpretations without making redundant copies of the text. In the Knora base ontology, any text value can have standoff nodes.

By representing standoff as RDF triples, Knora makes markup searchable across multiple text documents in a repository. For example, if a repository contains documents in which references to persons are indicated in standoff, it is straightforward to find all the documents mentioning a particular person. Knora’s standoff support is intended to make it possible to convert documents with embedded, hierarchical markup, such as TEI/XML, into RDF standoff and back again, with no data loss, thus bringing the benefits of RDF to existing TEI-encoded documents (cf. [7, p. 3]). While standoff nodes directly referring to positions in the text represent flat structures, the representation of XML-based markup or any hierarchical structures in standoff is enabled by introducing terminal and nonterminal nodes [9, 4]. Nonterminal nodes may represent hierarchical structures like lists: individual list elements refer to index positions in the text and point to a common nonterminal node representing their parent element (cf. [4, p. 229]). A purely flat approach (terminal nodes only) would mean that relations between single standoff nodes could not be handled at all by Knora. Taking the example of a list, the following embedded markup would not be represented correctly:

Issue 292:
Nonterminal
nodes do not
exist yet in
the ontology.

Listing 1: List in HTML

```
<ol>
  <li>first element</li>
  <li>second element</li>
  <li>third element</li>
</ol>
```

Applying a purely flat approach, only the index positions of the elements would be stored, but not the information that `` actually contains all the `` elements. This information would only implicitly be given because the `` spans the range of all the `` elements. By representing all the `` elements as terminal nodes and the `` as a nonterminal node, we are able to represent a truly hierarchical relation as standoff.⁶ The main benefit of using standoff nodes is that our approach may represent hierarchical relations (if

⁶Schmidt [7] differentiates between ‘standoff properties’ and ‘standoff markup’. The first term refers to an approach where attributes may freely overlap. The latter represents structures that comply with a context-free grammar such as an XML schema or a DTD (no overlap allowed). However, we think that hierarchies are not a restriction per se because there might be structures that are in fact hierarchical and should be represented as such also in standoff.

the relation we want to represent is in fact a hierarchy), but does not require markup information to be hierarchical due to syntactical restrictions.

In the Knora base ontology, a **TextValue** can have one or more standoff nodes. Each standoff node indicates the start and end positions of a substring in the text that has a particular attribute. The OWL class **kb:Standoff**, which is the base class of all standoff node classes, has these properties:

standoffHasAttribute (1) The name of the attribute.

standoffHasStart (1) The index of the first character in the text that has the attribute.

standoffHasEnd (1) The index of the last character in the text that has the attribute, plus 1.

The **Standoff** class is not used directly in RDF data; instead, its subclasses are used. A few subclasses are currently provided, and more will be added to support TEI semantics.

5.1 Subclasses of Standoff

5.1.1 StandoffVisualAttribute

Represents a formatting attribute such as *boldface* or **italics**. The value of **standoffHasAttribute** is the name of the formatting attribute, and can be any string. It is up to the text renderer to interpret these names.

5.1.2 StandoffHref

Indicates that a substring is associated with a resource on the Internet, i.e. a URL. It has this property:

standoffHasHref (1) An `xsd:anyURI` representing the URL.

5.1.3 StandoffLink

Indicates that a substring is associated with a Knora resource. For example, if a repository contains resources representing persons, a text could be marked up so that each time a person's name is mentioned, a **StandoffLink** connects the name to the Knora resource describing that person. It has this property:

standoffHasLink (1) The resource that the link points to.

How do we differentiate between truly hierarchical structures and structures that are just encoded as a hierarchy for a given XML file?

One of the design goals of the Knora ontology is to make it easy and efficient to find out which resources contain references to a given resource. Direct links are easier and more efficient to query than indirect links. Therefore, when a text value contains a resource reference in its standoff nodes, there must also be a direct link between the containing resource and the target resource, along with an RDF reification (a `kb:LinkValue`) describing the link, as discussed in Section 4. In this case, the link property is always `kb:hasStandoffLinkTo`, and the link value property (which points to the `LinkValue`) is always `kb:hasStandoffLinkToValue`.

The Knora API server automatically creates and updates direct links and reifications for standoff resource references when it creates and updates text values. To do this, it keeps track of the number of text values in each resource that contain at least one standoff reference to a given target resource. It stores this number as the reference count of the `LinkValue` (see Section 3.3.11) describing the direct link. Each time this number changes, it makes a new version of the `LinkValue`, with an updated reference count. When the reference count reaches zero, it removes the direct link and makes a new version of the `LinkValue`, marked with `kb:isDeleted`.

For example, if `data:R1` is a resource with a text value in which the resource `data:R2` is referenced, the repository could contain the following triples:

```
data:R1 ex:hasComment data:V1 .

data:V1 rdf:type kb:TextValue ;
  kb:valueHasString "This link is internal." ;
  kb:valueHasStandoff data:S01 .

data:S01 rdf:type kb:StandoffLink ;
  kb:standoffHasStart: 5 ;
  kb:standoffHasEnd: 9 ;
  kb:standoffHasLink data:R2 .

data:R1 kb:hasStandoffLinkTo data:R2 .
data:R1 kb:hasStandoffLinkToValue data:LV1 .

data:LV1 rdf:type kb:LinkValue ;
  rdf:subject data:R1 ;
  rdf:predicate kb:hasStandoffLinkTo ;
  rdf:object data:R2 ;
  kb:valueHasRefCount 1 .
```

Figure 6 illustrates the result.

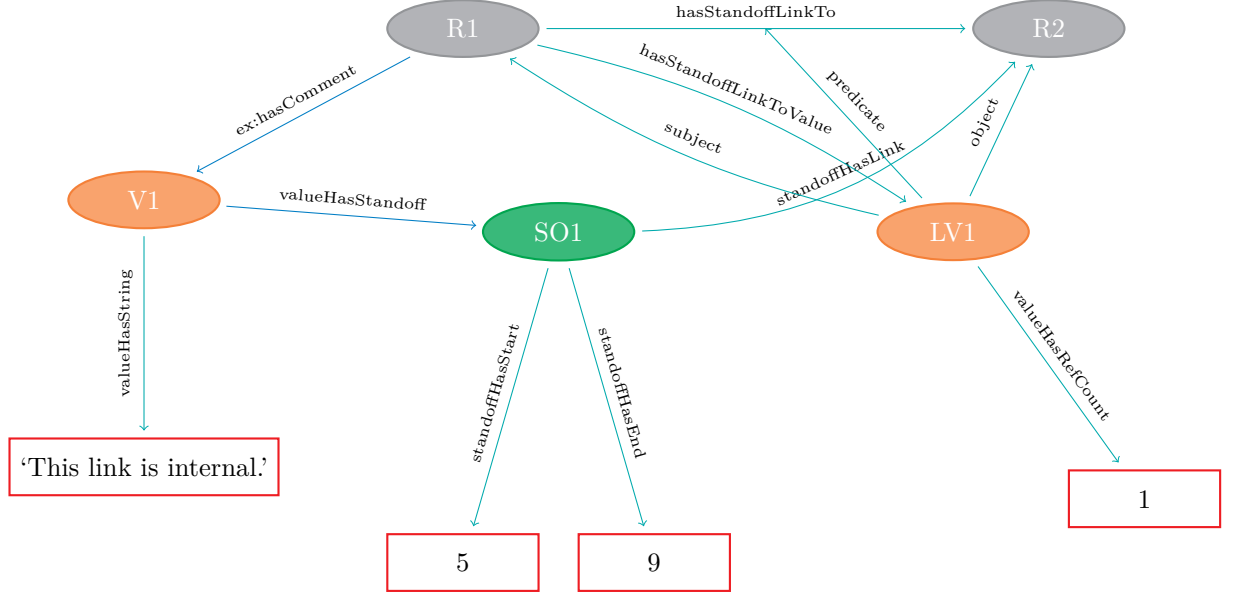


Figure 6: An RDF graph showing a link resulting from a reference to a resource in a standoff node.

6 Authorization

6.1 Users and Groups

Each Knora user is represented by an object belonging to the class `kb:User`, which is a subclass of `foaf:Person`, and has the following properties:

userid (1) A unique identifier that the user must provide when logging in.

password (1) A cryptographic hash of the user's password.

email (0-n) Email addresses belonging to the user.

isInProject (0-n) Projects that the user is a member of.

isInGroup (0-n) Project-specific groups that the user is a member of.

foaf:familyName (1) The user's family name.

foaf:givenName (1) The user's given name.

Knora's concept of access control is that an object (a resource or value) can grant permissions to groups of users (but not to individual users). There are four built-in groups:

UnknownUser Any user who has not logged into the Knora API server is automatically assigned to this group.

KnownUser Any user who has logged into the Knora API server is automatically assigned to this group.

ProjectMember When checking a user's permissions on an object, the user is automatically assigned to this group if she is a member of the project that the object belongs to.

Owner When checking a user's permissions on an object, the user is automatically assigned to this group if he is the owner of the object.

A project-specific ontology can define additional groups, which must belong to the OWL class `kb:UserGroup`.

6.2 Permissions

The owner of an object is always allowed to perform any operation on it. An object can grant the following permissions:

1. **hasRestrictedViewPermission:** Allows a restricted view of the object, e.g. a view of an image with a watermark.
2. **hasViewPermission:** Allows an unrestricted view of the object. Having view permission on a resource only affects the user's ability to view information about the resource other than its values. To view a value, she must have view permission on the value itself.
3. **hasModifyPermission:** For values, this permission allows a new version of a value to be created. For resources, this allows the user to create a new value (as opposed to a new version of an existing value), or to change information about the resource other than its values. When he wants to make a new version of a value, his permissions on the containing resource are not relevant. However, when he wants to change the target of a link, the old link must be deleted and a new one created, so he needs modify permission on the resource.
4. **hasDeletePermission:** Allows the item to be marked as deleted.

5. **hasChangeRightsPermission**: Allows the permissions granted by the object to be changed.

Each permission in the above list implies all lower-numbered permissions. A user's permission level on a particular object is calculated in the following way:

1. Make a list of the groups that the user belongs to, including **Owner** and/or **ProjectMember** if applicable.
2. If the user is the owner of the object, give her the highest level of permissions.
3. Otherwise, make a list of the permissions that she can obtain on the object, by iterating over the permissions that the object grants. For each permission, if she is in the specified group, add the specified permission to the list of permissions she can obtain.
4. From the resulting list, select the highest-level permission.
5. If the result is that she would have no permissions, give her whatever permission **UnknownUser** would have.

7 Open Questions

7.1 Extending Existing Resource Definitions

How should extensions of existing resources be handled? Project B extends a resource defined in the project A ontology, by adding new properties/values which are interesting for project B.

References

- [1] David Beckett et al. *RDF 1.1 Turtle: Terse RDF Triple Language*. Tech. rep. W3C, 2014. URL: <http://www.w3.org/TR/turtle/>.
- [2] Dan Brickley and R.V. Guha. *RDF Schema 1.1*. Tech. rep. W3C, 2014. URL: <http://www.w3.org/TR/2014/REC-rdf-schema-20140225/>.
- [3] Lou Burnard and Syd Bauman, eds. *Guidelines for Electronic Text Encoding and Interchange*. P5. Version 2.9.1. Last updated on 15th October 2015, revision 46ac023. Nov. 2007. URL: <http://www.tei-c.org/release/doc/tei-p5-doc/en/html/index.html>.

- [4] Christian Chiarcos. ‘POWLA: Modeling Linguistic Corpora in OWL/DL’. In: *The Semantic Web: Research and Applications, 9th Extended Semantic Web Conference 2012, Proceedings*. Ed. by Elena Simperl et al. Lecture Notes in Computer Science. Berlin, 2012, pp. 225–239.
- [5] M. Duerst and M. Suignard. *Internationalized Resource Identifiers (IRIs)*. Tech. rep. W3C, 2005. URL: <http://tools.ietf.org/html/rfc3987>.
- [6] Pascal Hitzler et al. *OWL 2 Web Ontology Language Primer (Second Edition)*. Tech. rep. W3C, 2012. URL: <http://www.w3.org/TR/2012/REC-owl2-primer-20121211/>.
- [7] Desmond Schmidt. ‘Using Standoff Properties for Marking-up Historical Documents in the Humanities’. In: *it – Information Technology* 58 (2016), p. 1. URL: <http://ecdosis.net/papers/schmidt.d.2016.pdf>.
- [8] Guus Schreiber and Yves Raimond. *RDF 1.1 Primer*. Tech. rep. W3C, 2014. URL: <http://www.w3.org/TR/2014/NOTE-rdf11-primer-20140624/>.
- [9] Amir Zeldes, Florian Zipser, and Arne Neumann. *PAULA XML Documentation: Format version 1.1*. Tech. rep. Universität Potsdam, Jan. 2013. URL: https://www.sfb632.uni-potsdam.de/images/doc/PAULA_P1.1.2013.1.21a.pdf.